

BoolQ with Pretrained Transformers

Yelin Zhang

Preprocessing

- Concat question and passage with [SEP]
- DeBERTA-v2 Tokenizer
 - SentencePiece subword tokenizer
 - SentencePiece embeddings
 - Compatible with DeBERTA-v3
- No other processing done, because SentencePiece expects raw text
- Padded by max length in every batch with [PAD]

Network architecture

Pretrained Transformer

- DeBERTA-v3 layers
 - Max input 512, Out 768
- ContextPooler
 - takes first hidden state [CLS]
 - In 768, Out 768

Classification Head


- Linear
 - In 768, Out 256
 - ReLu
- Linear
 - In 256, Out 1

Optimizer: AdamW

Loss: BCE with Logits

Experiments


- Learning rate in continuous range $1 - 1e-10$
 - potentially unlimited runs
- Completed 7 runs
 - due to long training times



Results

	Train	Valid	Test
Loss	0.6275	0.6839	0.6809
Accuracy	0.25	0.5950	0.6218

Best parameter: Learning rate 0.12



Comparison test metrics

	Majority class	Word embeddings	RNN	Transformer	Pre trained Transformer	
Accuracy	0.6218	0.6165	0.6397	0.6287	0.6218	

Interpretation

- LR over ~ 0.3 : pure minority class classifiers
- LR under ~ 0.3 : pure majority class classifiers
- Reasonable LR (probably at most ~ 0.01): mix of classifications
- Future: Sweep more restrictive range of learning rates